

MOHAN WAKCHAURE

+91 8805548797 | mohanwakchaure1@gmail.com | [LinkedIn](#) | [GitHub](#)
1007, Mangalam Signature, Moshi-Aalandi Road, Dudulgaon, Moshi, Pune, Maharashtra - 412105

PROFESSIONAL SUMMARY

AI & automation engineer with 3+ years building LLM-powered production tools that turn manual, compliance-sensitive operational work into reliable, governed automated workflows. Owns solutions end-to-end — requirement capture through build, deployment, monitoring, and iteration — spanning LLM tool/function calling and structured output, semantic retrieval over vector stores (RAG), multi-provider LLM routing with failover, human-in-the-loop approval steps, and AI guardrails / controls for automated actions. Strong in Python and secure REST / webhook / enterprise API integration (Microsoft Graph, Azure) with authentication, validation, and error handling, plus role-based access control, audit logging, PII-sensitive data handling, real-time monitoring, and cost-governed GCP deployment (Cloud Run, Cloud Build). Shipped AI document-processing, email-intelligence, and case-evaluation tools into daily production across 13+ teams.

EXPERIENCE

Operations Automation Specialist Jul 2026 – Present

Law Offices of Mary Kennedy Remote / United States

- Sole automation engineer for a US immigration law firm; own AI-tool development end-to-end from requirement capture through build, deployment, and iteration, coordinating directly with the CEO, front-desk lead, and prior developers.
- Re-architected the EB-1A case-evaluation tool (FastAPI, Python, Anthropic Claude, PostgreSQL/Cloud SQL) to remove a core compliance risk where an LLM was directly rendering legal eligibility determinations: replaced the AI-decides pipeline with a deterministic rules engine plus a mandatory human-review workflow, restricting AI to CV text extraction and advisory drafting only — a guardrail design with explicit human-in-the-loop control over automated actions.
- Built a legal-precedent retrieval layer with ChromaDB vector embeddings (semantic search) that grounds AI-drafted advisory suggestions and every human eligibility decision in cited precedent before acceptance (retrieval-augmented, citation-grounded).
- Engineered the admin platform — session-based authentication, role-based access control (owner/reviewer/front-desk), race-safe task-claiming, live/incremental scoring, and full audit logging (SQLAlchemy/Alembic) — with questionnaire auto-pre-fill from resume text and OpenAlex H-index enrichment for research-based candidates.
- Built a production bulk-import ETL pipeline migrating ~100 unstructured Excel lead records into the live system with deterministic email+phone dedup, a mandatory dry-run preview before any write, an isolated schema separating historical data from system-verified determinations, and inline source-correction with audit logging.
- Cut Cloud Run compute spend ~77% (verified against billing) by eliminating idle tagged revisions, and institutionalized a zero-downtime deploy/rollback pattern (zero-traffic tagged revision → health & log verification → cutover → tag cleanup) via Cloud Build with automated Artifact Registry image cleanup; rotated a production Claude API credential with zero downtime and closed a silent password-reset approval gap with targeted approver alerts.
- Built and deployed an AI email-intelligence system (Python, Flask, Linux/systemd) for the firm's front-desk team: scans incoming client email across 50+ M365 mailboxes, summarizes and classifies it by type/urgency via a multi-provider LLM pipeline (Groq/Llama with Mistral/Gemini failover) plus rule-based filtering, and delivers a timezone-aware twice-daily digest — from which supervisors view or forward the original message via stateless token-based action links (no login required), backed by a time-bound email cache.

Data Analyst / Automation Engineer Mar 2026 – Jul 2026

Printo Bengaluru

- Sole AI/automation point of contact across 13 teams; mapped manual, repetitive processes with product, operations, procurement, and finance stakeholders and translated them into automation solutions.
- Built a nightly Python ETL pipeline scanning 100,000+ SKU rows that triggers real-time low-stock alerts via WhatsApp Business API and Google Chat — replacing a manual daily check by 3 people with zero missed alerts since go-live.
- Built a Groq-powered LLM document-extraction tool (Streamlit, llama-3.1-8b-instant, pdfplumber, pytesseract) for Finance/AP that extracts fields from invoices and POs and auto-matches delivery challans, cutting manual data entry.
- Shipped 6 internal automation tools into daily production use, including approval-routing (human-in-the-loop) email workflows and a Firebase-backed material classification and inventory tool used by ~10 procurement and ops users; maintained ETL pipelines across Google Sheets, Firebase/Firestore, Supabase, and MySQL/PostgreSQL.

Fulfillment Analyst

May 2025 – Mar 2026

Meesho Bengaluru

- Automation point of contact for IT, Operations, Sales, HR, and Infrastructure teams; automated Jira incident and service-request workflows on Atlassian — auto-response bot, round-robin and shift-based ticket routing, and automated acknowledgement — reducing manual triage and assignment effort.
- Built Slack workflow automation for internal request handling; designed SQL, Excel, and Power BI trackers monitoring SLA breaches and compliance deviations, adopted into the procurement review process.

Machine Learning Intern Dec 2024 – May 2025

ExosphereX AI Pune

- Built and evaluated ML models for internal AI product features across preprocessing, feature engineering, training, and evaluation; integrated outputs into production workflows.

Software Engineer Aug 2024 – Nov 2024

Skylex Technology Pune

- Analysed campaign performance data (CTR, conversion) and built automated reporting pipelines that informed spend-reallocation decisions.

Machine Learning Intern Feb 2023 – Mar 2024

Indian Institute of Technology (IIT) Bombay Mumbai

- AI implementation: supported development and integration of AI/LLM features into internal tools and research workflows, contributing from prototype through integration.
- Contributed to LLM training / fine-tuning workflows — [dataset preparation, fine-tuning, and evaluation of models for X task].
- Built RPA workflows using Automation Anywhere to automate repetitive operational processes, reducing manual effort by [X hrs/week or X %].
- Developed operations-automation scripts (Python) that streamlined routine tasks and [outcome / metric].

SELECTED AUTOMATION & AI PROJECTS

Real-Time Inventory Monitoring & Alert Pipeline | Python, ETL, WhatsApp API, Google Chat, MySQL, Supabase — Production

- Nightly scan of 100,000+ SKU rows with threshold-based real-time alerts; replaced a manual daily check by 3 people with zero missed stockout alerts since go-live.

LLM Document Extraction Tool | Python, Groq (Llama-3.1-8b-instant), Streamlit, pdfplumber, pytesseract — Production

- Extracts structured fields from invoices and POs (digital and scanned) with automated challan-vs-PO matching; output to Excel via openpyxl.

Missed-Call Tracking & Callback Verification System | Google Apps Script, Vitel Telephony API, Google Sheets — Production

- Pulls and dedupes call-detail records from the Vitel telephony API and reconciles them against a live Google Sheet; a rule-based engine determines whether each call was resolved and sends staff twice-daily prioritized callback digests, with custom audit tooling that traced a recurring scheduled-trigger failure to a Google Sheets row-group-depth limit.

LLM Timekeeping Summarization | FastAPI, SQLite, Google Cloud Run, Google Gemini API — Built & Deployed

- LLM feature for a production timekeeping app that condenses employees' daily task logs into structured summaries (completed, pending, carried-over); degrades gracefully to a deterministic rule-based fallback on any API error or timeout, and enforces a guardrail keeping all billable-hour totals sourced from database aggregations, never model output.

SKILLS & TOOLS

AI & LLM Development: LLM API integration (Anthropic Claude, Groq/Llama, Gemini, Mistral, OpenAI-compatible), tool/function calling, structured output, prompt & context engineering, retrieval-augmented generation (RAG) & semantic search (ChromaDB vector embeddings), multi-provider LLM routing & failover, human-in-the-loop / approval workflows, AI guardrails & controls for automated actions, AI-based classification, document extraction

Automation & Bots: Workflow/ticket automation, incident & service-request automation (Jira/Atlassian), auto-response and routing bots, RPA (Automation Anywhere), process mapping, real-time monitoring & alerting, email automation

APIs & Integration: REST APIs, Webhooks, event-driven services, Microsoft Graph API, Azure app registration, Google Apps Script, Firebase/Firestore, Supabase

Languages & Data: Python (ETL, automation), SQL, PostgreSQL/Cloud SQL, MySQL, SQLite, SQLAlchemy, Alembic, Pydantic, Pandas, NumPy, scikit-learn

Cloud, DevOps & Governance: Google Cloud Platform (Cloud Run, Cloud Build, Cloud SQL, Artifact Registry, Compute Engine, Secret Manager, Cloud Storage, Cloud Logging), zero-downtime deployment & rollback, cost optimization, CI/CD, Git/GitHub, Linux / SSH / systemd, role-based access control, audit logging, PII-sensitive data handling, SharePoint / Microsoft Lists, Streamlit, FastAPI, Flask

Professional: Requirements gathering, stakeholder management, end-to-end ownership, troubleshooting & root-cause analysis, technical documentation

EDUCATION

M.Sc. — Artificial Intelligence and Machine Learning 2023 – 2025

MIT-ADT University, Pune, Maharashtra

B.Sc. — Computer Science (7.76/10) 2020 – 2023

SMBST College, Sangamner, Maharashtra